

MPI — An example

Salvatore Filippone

salvatore.filippone@uniroma2.it



The trapezoidal rule for computing a definite integral


If we have a sequence of points $a = x_0, x_1, x_2, \dots, x_{n-1}, b = x_n$, with constant $h = x_k - x_{k-1}$, and a function $f(x)$, we can approximate

$$\int_a^b f(x) \approx \sum_{i=1}^n \frac{1}{2} h [f(x_i) + f(x_{i-1})] = h \cdot \left\{ \frac{f(x_0)}{2} + \frac{f(x_n)}{2} + \sum_{i=1}^{n-1} f(x_i) \right\}.$$



An example problem

The trapezoidal rule for computing a definite integral





Definite integral — serial

```
#include <stdio.h>

float f(float x);

main(int argc, char* argv[]) {
    float integral;
    float a, b;
    int n;
    float h;
    int i;

    printf("Enter a, b and n\n");
    scanf("%f %f %d",&a,&b,&n);

    h=(b-a)/n;
    integral = (f(a)+f(b))/2.0;
    for (i=1; i<n; i++){
        x = a+i*h;
        integral = integral + f(x);
    }
    integral *= h;
    printf("With n=%d trapezoids we estimate integral from %f to %f: %f\n",
        n,a,b,integral);
}
```

Issues we need to figure out:

- Who gets the input;
- Who writes the output;
- How does everybody know what to do;
- How do you distribute the workload;
- What exactly is each process supposed to do;
- How is the team going to complete the process.

Issues we need to figure out:

- Who gets the input; **Process 0**
- Who writes the output; **Process 0**
- How does everybody know what to do; **Based on process rank; receive information from process 0**
- How do you distribute the workload; **As uniformly as possible;**
- What exactly is each process supposed to do; **Identify which portion it will compute, then compute**
- How is the team going to complete the process. **Send the results to process 0 which will assemble the final result**

```
#include <stdio.h>
#include <mpi.h>

float f(float x);
float Trap(float local_a, float local_b,
           int local_n, float h);

main(int argc, char* argv[]) {
    int p, my_rank;
    MPI_Status status;
    int i, tag;

    MPI_Init(&argc, &argv);

    MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);
    MPI_Comm_size(MPI_COMM_WORLD, &p);
```

Initial communication

```
tag = 0;
if (my_rank == 0) {
    printf("Enter a, b and n\n");
    scanf("%f %f %d",&a,&b,&n);
}
if (my_rank == 0) {
    for (i=1; i<p; i++) {
        MPI_Send(&a,1,MPI_FLOAT,i,tag,MPI_COMM_WORLD);
        MPI_Send(&b,1,MPI_FLOAT,i,tag,MPI_COMM_WORLD);
        MPI_Send(&n,1,MPI_INT,i,tag,MPI_COMM_WORLD);
    }
} else {
    MPI_Recv(&a,1,MPI_FLOAT,0,tag,MPI_COMM_WORLD);
    MPI_Recv(&b,1,MPI_FLOAT,0,tag,MPI_COMM_WORLD);
    MPI_Recv(&n,1,MPI_INT,0,tag,MPI_COMM_WORLD);
}
```


Local work setup & execution

```
float h, local_a, local_b, local_int;  
int local_n,  
  
h = (b-a)/n;  
local_n = n/p;  
  
// Start of local integration  
local_a = a+my_rank*local_n*h;  
local_b = local_a+local_n*h;  
  
local_int = Trap(local_a,local_b,local_n,h);
```

Collect data, print & close

```
float temp;

tag = 1;
if (my_rank == 0) {
    integral = local_int;
    for (i=1; i<p; i++){
        MPI_Recv(&temp,1,MPI_FLOAT,i,tag,
                 MPI_COMM_WORLD,&status);
        integral = integral + temp;
    }
    printf("With n=%d trapezoids we estimate integral",n);
    printf(" from %f to %f:  %f\n",a,b,integral);
} else {
    MPI_Send(&local_int,1,MPI_FLOAT,0,tag,
             MPI_COMM_WORLD,&status);
}

MPI_Finalize();
```

What is the main problem with the previous code?

What is the main problem with the previous code?

We have implicitly assumed that n is exactly divisible by p ! How do we fix this?

What is the main problem with the previous code?

We have implicitly assumed that n is exactly divisible by p ! How do we fix this?

```
int r;
r=n%p;
if (r == 0) {
    local_n = n/p;
    local_a = a+my_rank*local_n*h;
} else {
    if (my_rank < r) {
        local_n = n/p + 1;
        local_a = a+my_rank*local_n*h;
    } else {
        local_n = n/p+1;
        local_a = a+r*local_n*h;
        local_n = n/p;
        local_a = local_a + (my_rank-r)*local_n*h;
    }
    local_b = local_a+local_n*h;
}
```

Better yet

```
int r;
r=n%p;
if (my_rank < r) {
    local_n = n/p +1;
    local_a = a+my_rank*local_n*h;
} else {
    local_n = n/p+1;
    local_a = a+r*local_n*h;
    local_n = n/p;
    local_a = local_a + (my_rank-r)*local_n*h;
}
local_b = local_a+local_n*h;
```

(Why does it work?)

Anything else wrong with the previous code?

Anything else wrong with the previous code? Yes, the initial communication and the results collection

```
tag = 0;
if (my_rank == 0) {
    printf("Enter a, b and n\n");
    scanf("%f %f %d",&a,&b,&n);
}
if (my_rank == 0) {
    for (i=1; i<p; i++) {
        MPI_Send(&a,1,MPI_FLOAT,i,tag,MPI_COMM_WORLD);
        MPI_Send(&b,1,MPI_FLOAT,i,tag,MPI_COMM_WORLD);
        MPI_Send(&n,1,MPI_INT,i,tag,MPI_COMM_WORLD);
    }
} else {
    MPI_Recv(&a,1,MPI_FLOAT,0,tag,MPI_COMM_WORLD);
    MPI_Recv(&b,1,MPI_FLOAT,0,tag,MPI_COMM_WORLD);
    MPI_Recv(&n,1,MPI_INT,0,tag,MPI_COMM_WORLD);
}
```


Next topic: Collective Communications



Define and implement an MPI code that performs the Matrix-Vector product

$$y \leftarrow Ax;$$

consider as many variations as possible for the parallel data layout