

MPI Communicators, Groups, Topologies

Salvatore Filippone

salvatore.filippone@uniroma2.it



So far we have never really discussed one ubiquitous argument:

`MPI_COMM_WORLD`

Communicators

A communicator is used *both* to identify a group of processes *and* to separate messages

Why do you need a communicator?

Suppose you have a library, and you are invoking a function; suppose also that you are sending messages. How do you ensure your messages do not interfere with the library's messages?

The communicator is part of the *envelope*; as such, messages sent using one communicator cannot be received with another (even if it's an otherwise identical copy).

Communicators combine the concepts of

Group: a set of processes, each one of which has a *rank* from 0 to the size of the group (minus one);

Context: a partitioning of the communication space, such that messages in one context are separated from those in another.

For instance, all collective communication operations are executed in (and confined to) a certain group/context.

MPI_COMM_WORLD Predefined communicator, comprising all processes participating in the program execution; available immediately after a call to `MPI_Init`;

MPI_COMM_SELF Predefined communicator, comprising only the calling process;

MPI_COMM_NULL Predefined invalid communicator.

- `MPI_Comm_group(MPI_Comm comm, MPI_Group *group)`
returns the group associated with a communicator;
- `MPI_Comm_dup(MPI_Comm comm,
MPI_Comm *newcomm)`
Duplicates an existing communicator;
- `MPI_Group_incl(MPI_Group group, int n,
const int ranks[], MPI_Group *newgroup)`
Creates a new group from a subset of an existing one;

- `MPI_Comm_create(MPI_Comm comm, MPI_Group group, MPI_Comm *newcomm)`

Creates a new communicator from `group`, which is a subgroup of the group associated with `comm`; must be called by all processes in `comm`;

- `MPI_Comm_create_group(MPI_Comm comm, MPI_Group group, int tag, MPI_Comm *newcomm)`


Creates a new communicator from `group`, which is a subgroup of the group associated with `comm`; must be called by all processes in `group`;



Nearest-neighbour communications are very common: they are essential in many scientific applications. MPI defines some *Neighbour collective communications*


- `MPI_Neighbor_allgather(const void* sendbuf, int sendcount, MPI_Datatype sendtype, void* recvbuf, int recvcount, MPI_Datatype recvtype, MPI_Comm comm)`
- `MPI_Neighbor_alltoall(const void* sendbuf, int sendcount, MPI_Datatype sendtype, void* recvbuf, int recvcount, MPI_Datatype recvtype, MPI_Comm comm)`

V variants.





- `MPI_Cart_create(MPI_Comm comm_old, int ndims, const int dims[], const int periods[], int reorder, MPI_Comm *comm_cart)`
- `int MPI_Dims_create(int nnodes, int ndims, int dims[]),`





- `MPI_Graph_create(MPI_Comm comm_old, int nnodes, const int index[], const int edges[], int reorder, MPI_Comm *comm_graph)`
- `MPI_Dist_graph_create_adjacent(MPI_Comm comm_old, int indegree, const int sources[], const int sourceweights[], int outdegree, const int destinations[], const int destweights[], MPI_Info info, int reorder, MPI_Comm *comm_dist_graph)`



Exercises:

- implement a matrix-vector product with the matrix A distributed on a 2D cartesian topology;
- implement a matrix-matrix product with the matrices A , B and C distributed on a 2D cartesian topology;
- Implement a Jacobi sweep on a distributed cartesian mesh;