

Parallel Computing Systems and Applications

MPI — Extras

Salvatore Filippone

salvatore.filippone@uniroma2.it

Getting/setting environment parameters
Compile or run time:

```
#define MPI_VERSION 3  
#define MPI_SUBVERSION 1
```

```
int MPI_Get_version(int *version, int *subversion)
```

```
int MPI_Get_library_version(char *version, int *resultlen)
```

```
int MPI_Get_processor_name(char *name, int *resultlen)
```

Allocating (special-purpose) memory:

```
int MPI_Alloc_mem(MPI_Aint size, MPI_Info info, void *baseptr)
int MPI_Free_mem(void *base)
```

Example:

```
float (* f)[100][100];
/* no memory is allocated */
MPI_Alloc_mem(sizeof(float)*100*100, MPI_INFO_NULL, &f);
/* memory allocated */
...
(*f)[5][3] = 2.71;
...
MPI_Free_mem(f);
```

Error handling. Predefined error handlers:

```
MPI_ERRORS_ARE_FATAL
MPI_ERRORS_RETURN
```

User-defined error handlers:

```
typedef void MPI_Comm_errhandler_function(MPI_Comm *, int *, ...);
typedef void MPI_Win_errhandler_function(MPI_Win *, int *, ...);
typedef void MPI_File_errhandler_function(MPI_File *, int *, ...);

int MPI_Comm_create_errhandler(MPI_Comm_errhandler_function
                               *comm_errhandler_fn, MPI_Errhandler *errhandler)

int MPI_Comm_set_errhandler(MPI_Comm comm, MPI_Errhandler errhandler)
int MPI_Comm_get_errhandler(MPI_Comm comm, MPI_Errhandler *errhandler)
int MPI_Errhandler_free(MPI_Errhandler *errhandler)

int MPI_Win_create_errhandler(MPI_Win_errhandler_function
                              *win_errhandler_fn, MPI_Errhandler *errhandler)

int MPI_Win_set_errhandler(MPI_Win win, MPI_Errhandler errhandler)
int MPI_Win_get_errhandler(MPI_Win win, MPI_Errhandler *errhandler)
int MPI_File_create_errhandler(MPI_File_errhandler_function
                               *file_errhandler_fn, MPI_Errhandler *errhandler)
```

More on error handling

```
int MPI_Add_error_class(int *errorclass)
int MPI_Add_error_code(int errorclass, int *errorcode)
int MPI_Add_error_string(int errorcode, const char *string)

int MPI_Comm_call_errhandler(MPI_Comm comm, int errorcode)
int MPI_Win_call_errhandler(MPI_Win win, int errorcode)
int MPI_File_call_errhandler(MPI_File fh, int errorcode)
```

Timers:

```
double MPI_Wtime(void)
double MPI_Wtick(void)
```

Startup & friends

```
int MPI_Init(int *argc, char ***argv)
int MPI_Finalize(void)
int MPI_Initialized(int *flag)
int MPI_Finalized(int *flag)

int MPI_Abort(MPI_Comm comm, int errorcode)
```



It is possible to create (“spawn”) MPI processes dynamically. Given the enormous variety of execution environment it is impossible to give general advice, but it is possible to use the `MPI_COMM_WORLD` attribute

`MPI_UNIVERSE_SIZE`

to make informed decisions

```
int MPI_Comm_spawn(const char *command, char *argv[], int maxprocs,
    MPI_Info info, int root, MPI_Comm comm, MPI_Comm *intercomm,
    int array_of_errcodes[])
```

```
int MPI_Comm_get_parent(MPI_Comm *parent)
```



Dynamic Process Creation

```
/* manager */
#include "mpi.h"
int main(int argc, char *argv[])
{
    int world_size, universe_size, *universe_sizep, flag;
    MPI_Comm everyone; /* intercommunicator */
    char worker_program[100];
    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_COMM_WORLD, &world_size);
    if (world_size != 1)
        error("Top heavy with management");
    MPI_Comm_get_attr(MPI_COMM_WORLD, MPI_UNIVERSE_SIZE,
                      &universe_sizep, &flag);

    if (!flag) {
        printf("This MPI does not support UNIVERSE_SIZE. How many\n\
processes total?");
        scanf("%d", &universe_size);
    } else universe_size = *universe_sizep;
    choose_worker_program(worker_program);
    MPI_Comm_spawn(worker_program, MPI_ARGV_NULL, universe_size-1,
                  MPI_INFO_NULL, 0, MPI_COMM_SELF, &everyone, MPI_ERRCODES_IGNORE);
    MPI_Finalize();
    return 0;
}
```




Dynamic Process Creation

```
/* worker */
#include "mpi.h"
int main(int argc, char *argv[])
{
    int size;
    MPI_Comm parent;
    MPI_Init(&argc, &argv);
    MPI_Comm_get_parent(&parent);
    if (parent == MPI_COMM_NULL) error("No parent!");
    MPI_Comm_remote_size(parent, &size);
    if (size != 1) error("Something is wrong with the parent");
    /*
     * Parallel code here.
     * The manager is represented as the process with rank 0 in (the remote
     * group of) the parent communicator. If the workers need to communicate
     * among themselves, they can use MPI_COMM_WORLD.
     */
    MPI_Finalize();
    return 0;
}
```



From the MPI standard document

POSIX provides a model of a widely portable file system, but the portability and optimization needed for parallel I/O cannot be achieved with the POSIX interface.

The significant optimizations required for efficiency (e.g., grouping [47], collective buffering [7, 15, 48, 52, 58], and disk-directed I/O [43]) can only be implemented if the parallel I/O system provides a high-level interface supporting partitioning of file data among processes and a collective interface supporting complete transfers of global data structures between process memories and files. In addition, further efficiencies can be gained via support for asynchronous I/O, strided accesses, and control over physical file layout on storage devices (disks).

The I/O environment described in this chapter provides these facilities. Instead of defining I/O access modes to express the common patterns for accessing a shared file (broadcast, reduction, scatter, gather), we chose another approach in which data partitioning is expressed using derived datatypes. Compared to a limited set of predefined access patterns, this approach has the advantage of added flexibility and expressiveness.

The MPI I/O subsystem is powerful, but very complicated because it imposes a *structure* on the files, instead of using just byte sequence. Use it when necessary.



The MPI I/O interface

file An MPI file is an ordered collection of typed data items. A file is opened collectively by a group of processes. All collective I/O calls on a file are collective over this group.

displacement A file displacement is an absolute byte position relative to the beginning of a file.

etype An etype (elementary datatype) is the unit of data access and positioning.

filetype A filetype is the basis for partitioning a file among processes and defines a template for accessing the file

view A view defines the current set of data visible and accessible from an open file as an ordered set of etypes.



```
int MPI_File_open(MPI_Comm comm, const char *filename, int amode,
                  MPI_Info info, MPI_File *fh)
int MPI_File_close(MPI_File *fh)
int MPI_File_delete(const char *filename, MPI_Info info)
int MPI_File_set_size(MPI_File fh, MPI_Offset size)
int MPI_File_preallocate(MPI_File fh, MPI_Offset size)
int MPI_File_get_size(MPI_File fh, MPI_Offset *size)
int MPI_File_get_group(MPI_File fh, MPI_Group *group)
int MPI_File_get_amode(MPI_File fh, int *amode)
int MPI_File_set_view(MPI_File fh, MPI_Offset disp, MPI_Datatype etype,
                      MPI_Datatype filetype, const char *datarep, MPI_Info info)
int MPI_File_read_at(MPI_File fh, MPI_Offset offset, void *buf, int count,
                     MPI_Datatype datatype, MPI_Status *status)
int MPI_File_write_at(MPI_File fh, MPI_Offset offset, const void *buf,
                      int count, MPI_Datatype datatype, MPI_Status *status)
```



Used to implement new non-blocking operations

```
int MPI_Grequest_start(MPI_Grequest_query_function *query_fn,  
    MPI_Grequest_free_function *free_fn,  
    MPI_Grequest_cancel_function *cancel_fn, void *extra_state,  
    MPI_Request *request)  
typedef int MPI_Grequest_free_function(void *extra_state);  
typedef int MPI_Grequest_cancel_function(void *extra_state, int complete);  
int MPI_Grequest_complete(MPI_Request request)
```